

```

import tensorflow as tf
from tensorflow import keras
import matplotlib.pyplot as plt

fashion_mnist = keras.datasets.fashion_mnist
(X_train_full, y_train_full), (X_test, y_test) = fashion_mnist.load_data()

X_train_full.shape
X_train_full.dtype

X_valid, X_train = X_train_full[:5000]/255.0, X_train_full[5000:]/255.0
y_valid, y_train = y_train_full[:5000], y_train_full[5000:]
X_test = X_test / 255.0

```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz>  
 29515/29515 [=====] - 0s 0us/step  
 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz>  
 26421880/26421880 [=====] - 0s 0us/step  
 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz>  
 5148/5148 [=====] - 0s 0us/step  
 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz>  
 4422102/4422102 [=====] - 0s 0us/step

패션 MNIST 데이터셋을 적재한다.

```

[2] class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat', 'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']

[3] class_names[y_train[0]]

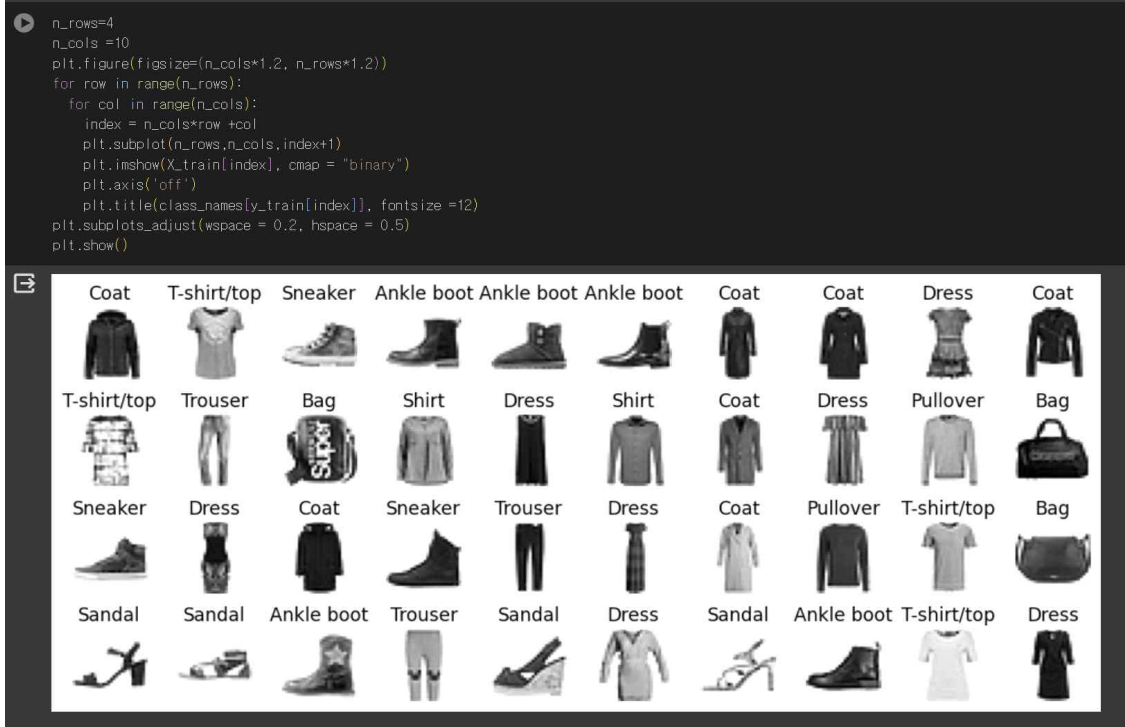
'Coat'

[4] plt.imshow(X_train[0], cmap="binary")
plt.axis('off')
plt.show()

```

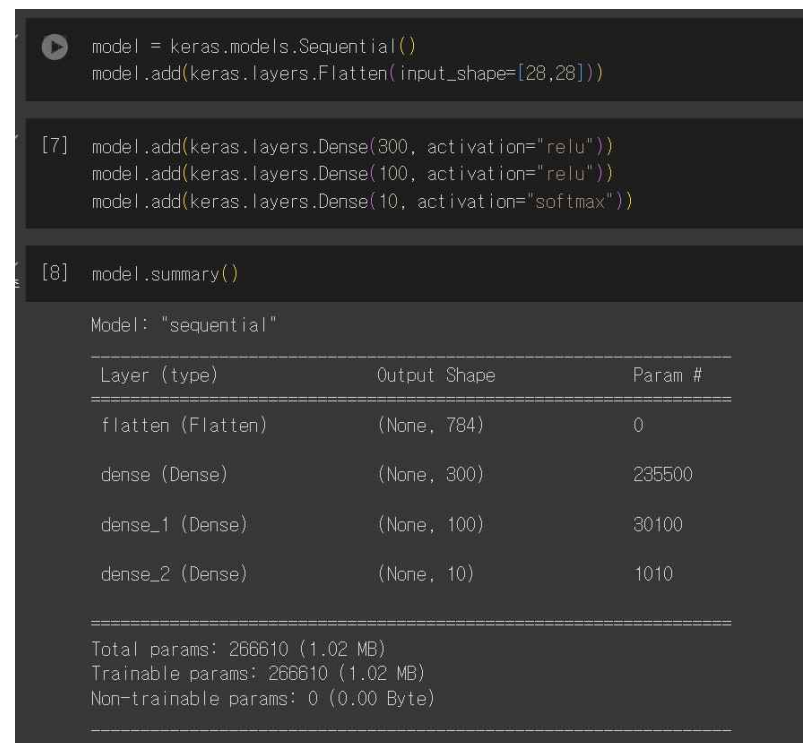


이미지를 확인해본다.

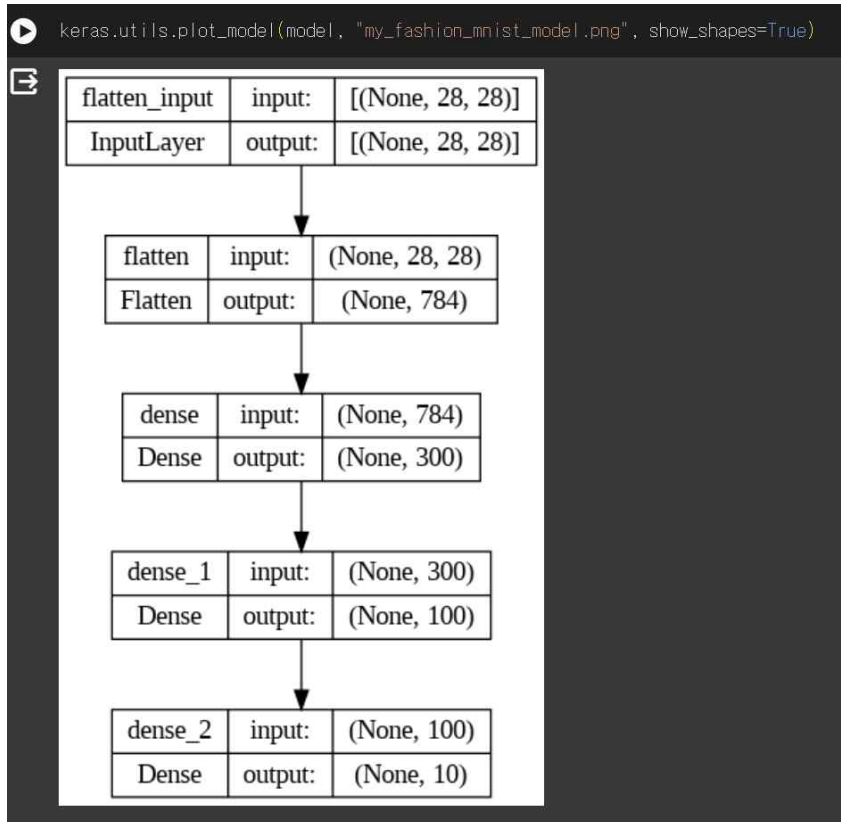


4행 10열로 훈련 이미지

40개를 출력한다.



Sequential API 로 모델을 만든 후 은닉층과 출력층을 추가하고 model.summary() 로 보여준다.



Keras 모델의 구조를 시각화한다.

```
[10] model.layers
for i in range(4):
    print(model.layers[i].name)

flatten
dense
dense_1
dense_2
```

Keras 모델에서 사용되는 모든 레이어의 목록을 반환한다.

```
[11] hidden1 = model.layers[1]

weights, biases = hidden1.get_weights()
```

모델의 두 번째 레이어를 추출하여 hidden1 변수에 저장한 후 hidden1에서 가중치와 편향을 추출하여 weights 와 biases 변수에 각각 저장한다.

```
[13] weights.shape
(784, 300)

[14] biases.shape
(300,)
```

가중치와 편향의 형태를 각각 출력한다.

```
[15] weights

array([[ 0.02196094, -0.01604879,  0.05912909, ...,  0.07435071,
         0.04544255,  0.00582179],
       [-0.01522749,  0.04079878, -0.06916433, ...,  0.00088407,
         0.01567963,  0.03441385],
       [ 0.04541522,  0.05755201, -0.03514511, ...,  0.0281071 ,
        -0.01258096,  0.02456383],
       ...,
       [-0.07360154,  0.02332964, -0.00932786, ..., -0.0287984 ,
         0.02989011,  0.02534214],
       [ 0.06294772,  0.04587 , -0.0346186 , ..., -0.01705816,
        -0.05546772,  0.03612237],
       [-0.03427893,  0.06364237, -0.04471117, ..., -0.06345218,
         0.07302211, -0.06511087]], dtype=float32)
```

가중치를 출력한다.

```
[16] model.compile(loss="sparse_categorical_crossentropy",
                  optimizer = "sgd",
                  metrics = ["accuracy"])
```

모델을 컴파일한다.

```
history = model.fit(X_train,y_train,epochs = 30,
                    validation_data = (X_valid, y_valid))

Epoch 1/30
1719/1719 [=====] - 17s 9ms/step - loss: 0.7189 - accuracy: 0.7649 - val_loss: 0.5134 - val_accuracy: 0.8280
Epoch 2/30
1719/1719 [=====] - 7s 4ms/step - loss: 0.4887 - accuracy: 0.8288 - val_loss: 0.4820 - val_accuracy: 0.8328
Epoch 3/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.4463 - accuracy: 0.8435 - val_loss: 0.4223 - val_accuracy: 0.8530
Epoch 4/30
1719/1719 [=====] - 7s 4ms/step - loss: 0.4199 - accuracy: 0.8511 - val_loss: 0.4015 - val_accuracy: 0.8668
Epoch 5/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.3982 - accuracy: 0.8602 - val_loss: 0.3846 - val_accuracy: 0.8694
Epoch 6/30
1719/1719 [=====] - 7s 4ms/step - loss: 0.3811 - accuracy: 0.8645 - val_loss: 0.3711 - val_accuracy: 0.8730
Epoch 7/30
1719/1719 [=====] - 6s 4ms/step - loss: 0.3679 - accuracy: 0.8695 - val_loss: 0.3684 - val_accuracy: 0.8722
Epoch 8/30
1719/1719 [=====] - 8s 5ms/step - loss: 0.3559 - accuracy: 0.8734 - val_loss: 0.3559 - val_accuracy: 0.8738
Epoch 9/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.3451 - accuracy: 0.8771 - val_loss: 0.3592 - val_accuracy: 0.8750
Epoch 10/30
1719/1719 [=====] - 6s 4ms/step - loss: 0.3360 - accuracy: 0.8811 - val_loss: 0.3658 - val_accuracy: 0.8692
Epoch 11/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.3275 - accuracy: 0.8837 - val_loss: 0.3689 - val_accuracy: 0.8636
Epoch 12/30
1719/1719 [=====] - 7s 4ms/step - loss: 0.3197 - accuracy: 0.8853 - val_loss: 0.3320 - val_accuracy: 0.8822
Epoch 13/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.3125 - accuracy: 0.8880 - val_loss: 0.3260 - val_accuracy: 0.8816
Epoch 14/30
1719/1719 [=====] - 7s 4ms/step - loss: 0.3059 - accuracy: 0.8897 - val_loss: 0.3266 - val_accuracy: 0.8844
Epoch 15/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2984 - accuracy: 0.8927 - val_loss: 0.3392 - val_accuracy: 0.8802
```

```
Epoch 16/30
1719/1719 [=====] - 8s 4ms/step - loss: 0.2927 - accuracy: 0.8941 - val_loss: 0.3259 - val_accuracy: 0.8790
Epoch 17/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2868 - accuracy: 0.8969 - val_loss: 0.3197 - val_accuracy: 0.8818
Epoch 18/30
1719/1719 [=====] - 7s 4ms/step - loss: 0.2808 - accuracy: 0.8992 - val_loss: 0.3330 - val_accuracy: 0.8772
Epoch 19/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2762 - accuracy: 0.9007 - val_loss: 0.3089 - val_accuracy: 0.8888
Epoch 20/30
1719/1719 [=====] - 6s 4ms/step - loss: 0.2701 - accuracy: 0.9033 - val_loss: 0.3335 - val_accuracy: 0.8804
Epoch 21/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2660 - accuracy: 0.9039 - val_loss: 0.3201 - val_accuracy: 0.8832
Epoch 22/30
1719/1719 [=====] - 6s 4ms/step - loss: 0.2600 - accuracy: 0.9058 - val_loss: 0.2997 - val_accuracy: 0.8898
Epoch 23/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.2567 - accuracy: 0.9076 - val_loss: 0.3170 - val_accuracy: 0.8914
Epoch 24/30
1719/1719 [=====] - 7s 4ms/step - loss: 0.2520 - accuracy: 0.9089 - val_loss: 0.3133 - val_accuracy: 0.8876
Epoch 25/30
1719/1719 [=====] - 7s 4ms/step - loss: 0.2474 - accuracy: 0.9098 - val_loss: 0.2996 - val_accuracy: 0.8914
Epoch 26/30
1719/1719 [=====] - 7s 4ms/step - loss: 0.2423 - accuracy: 0.9134 - val_loss: 0.2964 - val_accuracy: 0.8944
Epoch 27/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2382 - accuracy: 0.9133 - val_loss: 0.3140 - val_accuracy: 0.8908
Epoch 28/30
1719/1719 [=====] - 7s 4ms/step - loss: 0.2350 - accuracy: 0.9157 - val_loss: 0.3009 - val_accuracy: 0.8920
Epoch 29/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2312 - accuracy: 0.9166 - val_loss: 0.3033 - val_accuracy: 0.8908
Epoch 30/30
1719/1719 [=====] - 7s 4ms/step - loss: 0.2280 - accuracy: 0.9182 - val_loss: 0.2924 - val_accuracy: 0.8912
```

주어진 학습 데이터와 검증 데이터에 대해 모델을 30번의 epoch 동안 학습시킨다.

최종 훈련 세트 정확도는 91.82% 가 나왔고, 최종 검증 세트 정확도는 89.12% 가 나왔다.

```
help(model.fit)

Help on method fit in module keras.src.engine.training:

fit(x=None, y=None, batch_size=None, epochs=1, verbose='auto', callbacks=None, validation_split=0.0, validation_data=None)
Trains the model for a fixed number of epochs (dataset iterations).

Args:
  x: Input data. It could be:
    - A Numpy array (or array-like), or a list of arrays
      (in case the model has multiple inputs).
    - A TensorFlow tensor, or a list of tensors
      (in case the model has multiple inputs).
    - A dict mapping input names to the corresponding array/tensors,
      if the model has named inputs.
    - A `tf.data` dataset. Should return a tuple
      of either `(inputs, targets)` or
      `(inputs, targets, sample_weights)`.
    - A generator or `keras.utils.Sequence` returning `(inputs,
      targets)` or `(inputs, targets, sample_weights)`.
    - A `tf.keras.utils.experimental.DatasetCreator`, which wraps a
      callable that takes a single argument of type
      `tf.distribute.InputContext`, and returns a `tf.data.Dataset`.
```

Keras의 fit 메서드에 대한 도움말을 확인한다.

```
[19] history.params

{'verbose': 1, 'epochs': 30, 'steps': 1719}

[20] print(history.epoch)

[0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29]

[21] history.history.keys()

dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

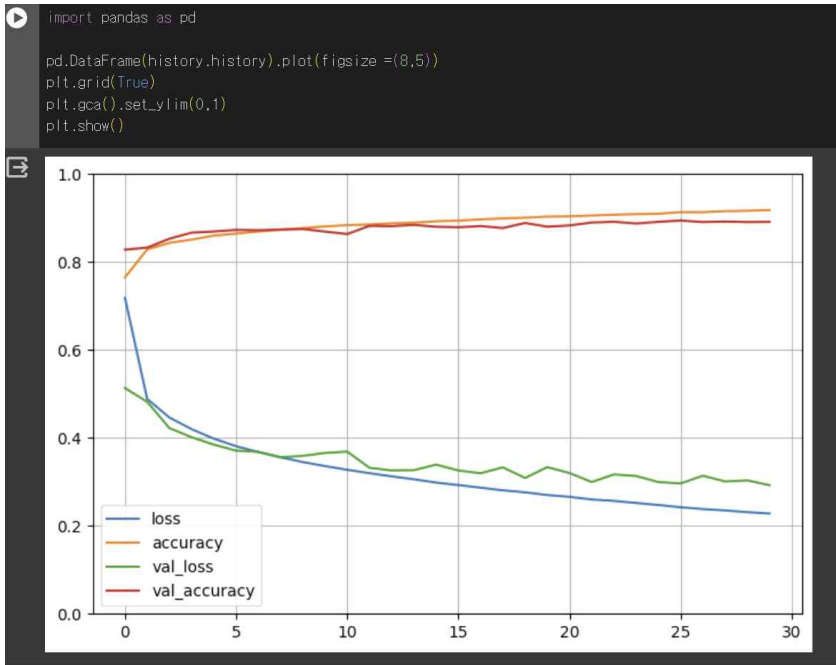
history.params 코드는 학습에 사용된 매개 변수들의 값들을 나타낸다.

그 후 각 epoch의 인덱스를 출력하고 학습 이력에 저장된 지표들의 이름을 출력한다.

```
history.history['accuracy']

[0.7648909091949463,
 0.8287818431854248,
 0.8435090780258179,
 0.8510727286338806,
 0.8602181673049927,
 0.8645272850990295,
 0.869490921497345,
 0.8733817934989929,
 0.8771091103553772,
 0.8811091184616089,
 0.8837454319000244,
 0.8852909207344055,
 0.8880000114440918,
 0.8896727561950684,
 0.8927090764045715,
 0.8941272497177124,
 0.8969454765319824,
 0.8992000222206116,
 0.9007090926170349,
 0.9032545685768127,
 0.9039090871810913,
 0.9058363437652588,
 0.9075999855995178,
 0.9089272618293762,
 0.9097636342048645,
 0.913418173789978,
 0.913345456123352,
 0.9157272577285767,
 0.9165818095207214,
 0.918218195438385]
```

history 객체에서 학습 중에 각 epoch에서의 정확도를 나타내는 리스트 또는 배열을 가져온다.



pandas를 사용하여 history 객체에 저장된 지표들을 데이터 프레임으로 변환한 후 시각화한다.

```
model.evaluate(X_test,y_test)
```

```
313/313 [=====] - 1s 2ms/step - loss: 0.3289 - accuracy: 0.8822
[0.3289021849632263, 0.8822000026702881]
```

테스트 세트로 최종 일반화 오차를 추정한다.


최종 테스트 세트 정확도는 88.22%이다.

```
X_new = X_test[:3]
y_proba = model.predict(X_new)
y_proba.round(2)
```

```
1/1 [=====] - 0s 104ms/step
array([[0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0.01, 0. , 0.98],
       [0. , 0. , 1. , 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. ],
       [0. , 1. , 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. ]],
      dtype=float32)
```

모델을 사용하여 새로운 데이터에 대한 예측을 수행하고, 각 클래스에 속할 확률을 출력한다.

```
[26] plt.figure(figsize=(7.2,2.4))
for index,image in enumerate(X_new):
    plt.subplot(1,3,index+1)
    plt.imshow(image, cmap = "binary")
    plt.axis('off')
    plt.title(class_names[y_test[index]],fontsize=12)
plt.subplots_adjust(wspace=0.2, hspace=0.5)
plt.show()
```

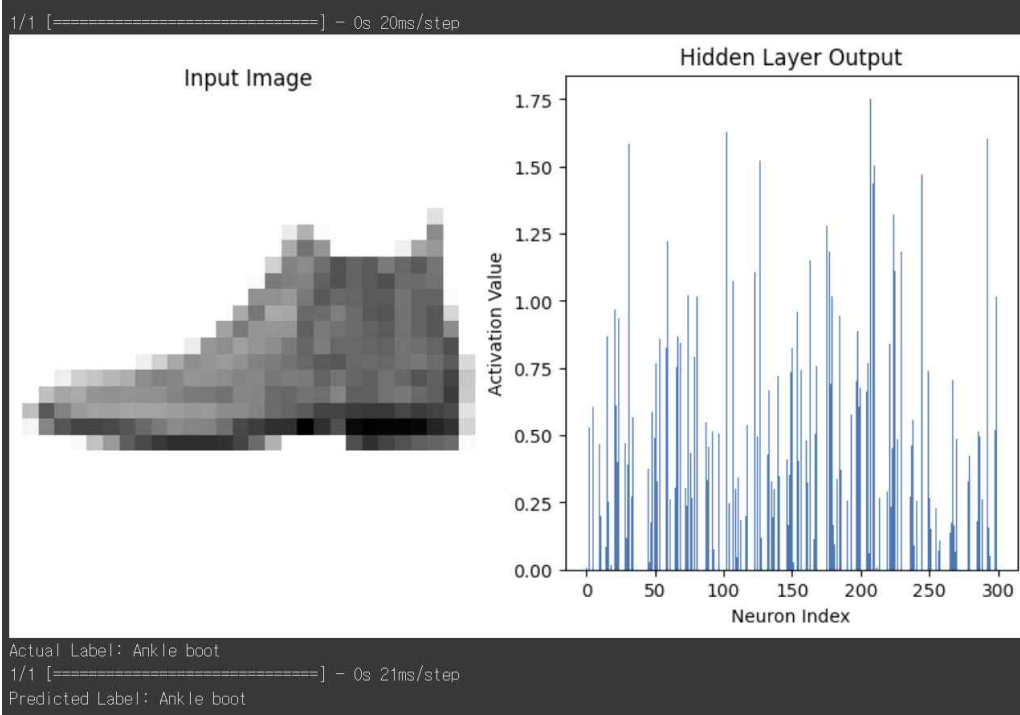


예측한 그림을 확인한다.

```
hidden_layer_output_model = keras.models.Model(inputs = model.input, outputs = model.layers[1].output)
```

```
[30] def visualize_hidden_layer_output(image_index):  
    image = X_test[image_index]  
    hidden_layer_output = hidden_layer_output_model.predict(image.reshape(1,28,28))  
    plt.figure(figsize = (10,5))  
    plt.subplot(1,2,1)  
    plt.imshow(image, cmap = "binary")  
    plt.title("Input Image")  
    plt.axis('off')  
  
    plt.subplot(1,2,2)  
    plt.bar(range(len(hidden_layer_output.flatten())), hidden_layer_output.flatten())  
    plt.title("Hidden Layer Output")  
    plt.xlabel("Neuron Index")  
    plt.ylabel("Activation Value")  
  
    plt.show()  
  
    print("Actual Label:", class_names[y_test[image_index]])  
    predicted_class_index = tf.argmax(model.predict(image.reshape(1, 28, 28)), axis=1).numpy()[0]  
    print("Predicted Label:", class_names[predicted_class_index])  
    image_index = 0  
    visualize_hidden_layer_output(image_index)
```

은닉층의 출력을 시각화하고, 해당 이미지의 실제 레이블을 함께 표시하는 코드이다.



코드 수행 결과이다.